

Cyborg groups enhance face recognition in crowded environments

Davide Valeriani^{1,2*}, Riccardo Poli²

1 Brain Computer Interfaces and Neural Engineering Laboratory, School of Computer Science and Electronic Engineering, University of Essex, Colchester, UK.

2 Department of Otolaryngology, Massachusetts Eye and Ear, Harvard Medical School, Boston, MA, USA.

* davide.valeriani@gmail.com

Abstract

Recognizing a person in a crowded environment is a challenging, yet critical, visual-search task for both humans and machine-vision algorithms. This paper explores the possibility of combining a residual neural network (ResNet), brain-computer interfaces (BCIs) and human participants to create “cyborgs” that improve decision making. Human participants and a ResNet undertook the same face-recognition experiment. BCIs were used to decode the decision confidence of humans from their EEG signals. Different types of cyborg groups were created, including either only humans (with or without the BCI) or groups of humans and the ResNet. Cyborg groups decisions were obtained weighing individual decisions by confidence estimates. Results show that groups of cyborgs are significantly more accurate (up to 35%) than the ResNet, the average participant, and equally-sized groups of humans not assisted by technology. These results suggest that melding humans, BCI, and machine-vision technology could significantly improve decision-making in realistic scenarios.

Introduction

Visual search is the process of looking for an item of interest in a scene. Everyone engages in visual search many times every day [1]. For example, when people look for their keys in a drawer, or when a bear hunts for fish in a river. Evolution has made animals very effective at searching in natural environments [2]. Yet, our performance is far from ideal [1,3]. For example, when the item of interest has a size that is inconsistent with the rest of the scene, we are more likely to miss it [4]. Phenomena such as inattentional blindness [5] and illusory conjunctions [6] lead us to miss objects or perceive them with wrong features.

These difficulties are present also when searching for a target face in pictures of crowded environments, a critical task in security and surveillance [7]. This challenging visual-search task requires humans to scan an image, detect faces, and compare them with the memory representation we have of a target person [8]. Even if our brain has dedicated regions to process face information [9], humans still find this task demanding and difficult, even when highly-skilled [10].

Searching for a target face in a crowded environment is a challenging task also for fully-automated computer-vision systems [11], albeit not for the same reasons as humans. Firstly, a module of the system needs to scan the input image and extract all the faces with their location (*face detection*) [12]. Secondly, another module has to compare each extracted face against the image of the target face and return a degree of match (*face identification*) [12]. If the match is above a certain threshold, the image is classified as containing the target. Both face detection and face identification unavoidably produce false alarms and misses. Despite these difficulties, research has shown that deep neural networks do not suffer from some limitations of the human visual system [4]. In principle this means that computer-vision systems could become more accurate than the average human [13], especially in tasks where visual information is presented for a brief amount of time. Unlike humans, the performance of computer-vision systems does not degrade over time (e.g., due to fatigue). However, it usually degrades way below that of humans when moving from constrained to realistic environments [14, 15].

To partially address the limitations of the human visual system, people can work together to search a scene. When making perceptual decisions, groups are generally

more accurate than their members [16]. This “wisdom of crowds” [17] has been found in
32 a variety of tasks, ranging from estimating uncertain quantities [18] to medical
33 diagnosis [19] and face matching [20]. However, there are circumstances in which groups
34 decisions are *worse* than individual ones. For example, when the opinions of group
35 members are too correlated [16], when a strong leader dominates the discussion [21], or
36 when group interactions are not effective [22, 23].
37

The strategy used to aggregate individual opinions critically influences performance
38 in group decision making. In the case of binary judgements, standard majority is the
39 most straightforward strategy to use. However, in many circumstances it is
40 suboptimal [24]. A well-chosen strategy could let groups make correct decisions even if
41 the majority (but not all) of the members are wrong [25]. For instance, this could be
42 achieved using a strategy consisting in weighing opinions on the basis of individual
43 confidence [26], which can be thought to approximate the probability of a decision being
44 correct [27]. Nevertheless, an accurate assessment of the reliability of the decision is
45 difficult to obtain. Humans may be miscalibrated in estimating their own
46 confidence [28], e.g., people may report high values of confidence when they made an
47 incorrect decision. This miscalibration is even accentuated by experience [29] and social
48 influence [30].
49

In principle, decision confidence could be more objectively assessed via means other
50 than self-reporting. There is ample evidence that physiological correlates of decision
51 confidence exist. For instance, several brain-activity patterns as recorded via
52 electroencephalography (EEG) are known to provide information about decision
53 confidence (e.g., the P300 event-related potential (ERP) and the error related
54 negativity [31, 32]), and visual search performance (e.g., the N2pc [33] and N1 [34]
55 ERPs). Moreover, it has been known for a long time that response times (RTs)
56 negatively correlate with the confidence in a decision [35–37]. An important question is,
57 however, whether these findings can be converted into practical systems to aid decision
58 making.
59

In recent research, we developed a series of *hybrid Brain-Computer Interfaces*
60 (hBCIs) to test whether it is indeed possible to decode the decision confidence of single
61 users from their EEG signals, RTs and other physiological measurements. In all cases,
62 the hBCI relies heavily on machine-learning technology, which is used to specialise the
63

interface to each particular user and to learn to predict the probability that, in a certain trial, the user made a correct decision. The confidence estimates provided by the hBCI were then used to weigh individual opinions and obtain group decisions. That is, our hBCI did not make automated decisions, it only produced weights for integrating behavioural decisions provided by the operators. These hBCI-assisted group decisions were then compared to decisions obtained by other forms of opinion aggregation. After initial tests of our system with a simple visual-matching task [38], which showed that hBCI-assisted groups were significantly more accurate than equally-sized groups based on standard majority, we progressed to much more realistic situations. For instance, we have shown that similar accuracy improvements could be obtained in visual search in crowded environments [39, 40], where participants were presented with cluttered scenes for 250 ms and had to decide whether a target object was present. The confidence estimated by our hBCI was also a more reliable predictor of correctness than the confidence reported by the participants themselves [40]. More recently, we have used our hBCI to assist humans searching for a target face in a crowded, indoor environment [41], which is of critical importance in many domains and is the application area considered also in this article. Once again, even in this real-world situation, groups of humans assisted by our hBCI were significantly more accurate than traditional groups based on standard majority or reported confidence [41]. While this is encouraging, with or without the assistance of an hBCI, even trained operators find this task hard and very fatiguing.

Given advantages and limitations of both humans (alone or in a group, and with or without the assistance of an hBCI) and computer-vision systems, one may wonder whether we could usefully *combine* these together. At least in principle, within such “cyborgs”, automated algorithms could compensate for the weaknesses of humans and *vice versa*, thereby producing more reliable and accurate decisions. For example, in a situation where most users in a group missed a target face (e.g., because the target appeared in an unusual/unsuspected location), both traditional and hBCI-assisted groups would likely make a wrong decision. Conversely, a machine-vision algorithm might still identify the target with confidence. So, if it was included as a member in a group of humans, the machine-vision algorithm could potentially swing the group decision towards a correct decision.

This study explores the benefits and drawbacks of this idea. Prior research has
96 shown that aggregating decisions from a superrecognizer and a deep convolutional
97 neural network in a face identification task led to significant more accurate decisions
98 than two superrecognitioners [42]. We used a state-of-the-art residual deep neural network
99 (ResNet), our BCI estimating the decision confidence using only neural features, and
100 human participants to create different types of cyborgs and investigate whether they
101 can improve decision making and why. We evaluated these approaches in the context of
102 the same face-recognition experiment as in [41], where participants had to *search* an
103 image of a crowded, indoor environment and decide whether or not it contained a target
104 face, a more challenging task than the face identification task used in [42].
105

Methods

Ethics statement.

This research received UK Ministry of Defence Research Ethics Committee (MoDREC,
108 ref. 520/MODREC/14) and University of Essex ethical approval in July 2014. All
109 participants signed an informed consent form before taking part in the experiment.
110

Participants.

Ten healthy participants (37.8 ± 4.8 years old, seven females, all right-handed) with
112 normal or corrected-to-normal vision and no reported history of epilepsy took part in
113 the experiment. Each participant was paid GBP 16 for taking part in the experiment
114 plus an additional amount of up to GBP 4 proportional to the participant's percentage
115 of correct decisions. The latter was used to encourage participants to focus on the task
116 and achieve maximum performance.
117

Stimuli and procedure.

Participants were comfortably seated at about 80 cm from a LCD screen (refresh rate
119 60 Hz) in an experimental room with normal white illumination (Fig 1(a)).
120

The experiment consisted of six blocks of 48 trials. At the beginning of each block,
121 participants were shown a cropped face image of the person to be considered as a target
122

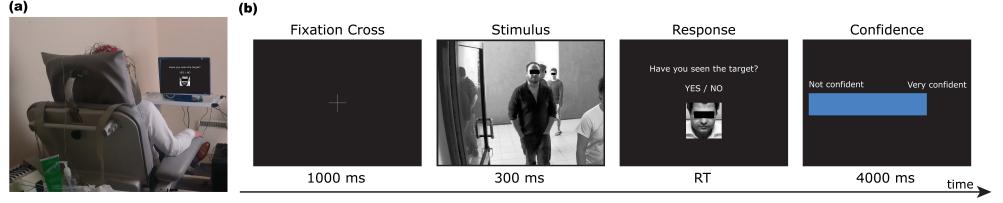


Fig 1. Experimental protocol. **(a)** Participant during one recording session. **(b)** Sequence of stimuli presented in a trial. Images representing people have been blurred for publication requirements.

for that block and were asked to memorise it before starting the presentation of the trials in the block. Each trial (Fig 1(b)) started with the presentation of a fixation cross for 1 s, followed by an image of a crowded scene presented for 300 ms in fullscreen, subtending approximately 14.4 degrees horizontally and 11.0 degrees vertically. Then, participants were presented with a display showing (again) the image of the target face and were asked to decide whether or not the target person was present in the scene, by clicking the left or the right mouse buttons, respectively. All volunteers were instructed to make decisions as quickly as possible. Finally, participants were asked to indicate, within 4 s, their subjective degree of confidence in that decision (0–100%) using the mouse wheel to vary the selected value in steps of 10%. Subjective confidence values were not used in this study as in preliminary analysis [41] we found that these confidence estimates were less accurate and robust than the confidence estimated by the BCI from the brain signals.

The images used as stimuli were obtained from the sequences P2L_S5 and P2E_S5 of the ChokePoint video dataset [43]. The two sequences contained video streams of 700+ frames, representing 29 people (six females) walking indoor through two different portals. For each sequence, three videos were recorded from cameras positioned at the top-left (L), top-center (C), and top-right (R) of each portal, respectively. From each sequence, we selected one person as a “target”. Then, we randomly selected 48 frames out of the 700+ available, 12 where the “target” person was visible in the scene and 36 where it was not. Since each frame was represented by three images taken from the three viewpoints, we selected a total of $2 \times 48 \times 3 = 288$ images. Each image, containing between 2 and 11 people, was converted to greyscale and had its histogram equalised.

In each block of the experiment, the images selected from a given sequence (1 or 2)

and viewpoint (L, C or R) were presented. All six possible combinations of sequences
147 and viewpoints have been tested. The order of presentation of the images in a particular
148 block was the same for all participants. However, the order of presentation of the blocks
149 was pseudo-randomly chosen for each participant, for counterbalancing reasons. In each
150 block 25% of the images contained the target.
151

Briefing, preparation and task familiarisation (via two blocks of 10 trials each) took
152 about 45 minutes, while the experiment took roughly 25 minutes. The images used as
153 stimuli during familiarisation were pseudo-randomly selected for each participant from
154 the dataset used in the experiment.
155

Data recording and BCI confidence estimation. 156

Neural data were recorded at 2048 Hz from 64 electrode locations with a BioSemi
157 ActiveTwo EEG system. Each channel was referenced to the average voltage recorded
158 from the electrodes placed on each earlobe and band-pass filtered between 0.15 and
159 40 Hz. Artefacts caused by ocular movements were removed with a standard
160 correlation-based subtraction algorithm. From each trial, response-locked epochs
161 starting 1 s before the user's response and lasting 1.5 s were extracted from the EEG
162 data, baseline corrected with the average voltage recorded in the 200 ms before the
163 stimulus onset, and downsampled to 32 Hz. Each epoch was then associated to the class
164 "correct" (confident) or "incorrect" (not confident) depending on whether the decision
165 made by the participant in that trial was correct or not, respectively.
166

Common Spatial Pattern (CSP) [44] filtering was used to transform the data in a
167 subspace where the variance between the "correct" and "incorrect" classes was
168 maximum. For each participant, we computed the CSP transformation matrix using the
169 epochs in the training set, and used such a matrix to transform the data in the test set.
170 The logarithm of the variances of the first and the last rows were computed and used as
171 neural features.
172

For each participant p , we trained a logistic regression classifier with L2 penalty to
173 predict the decision confidence w_p , i.e., the probability of a particular decision being
174 correct. Hence, each participant had his/her own BCI able to assess the objective
175 decision confidence of the decision maker. We used 8-fold cross-validation to ensure that
176

the results were not unfairly affected by overfitting. In each fold, 252 trials were used as
177 training set and the remaining 36 as test set, to evaluate the groups' performance.
178

ResNet architecture and training.

179

We used a residual neural network (ResNet) (see `face_recognition` Python library
180 v1.2.3 available at https://github.com/ageitgey/face_recognition) to decide
181 whether or not the target person was present in each image used as a stimulus in our
182 experiment. The ResNet was formed of 29 convolutional layers (similarly to the
183 ResNet-34 described by He et al. [45]), and had previously been trained on about 3
184 million faces derived from a number of datasets and tested on the “Labeled Faces in the
185 Wild” benchmark [46], where it achieved 99.38% accuracy. We ran our analysis using
186 Python 3.6.7.
187

In each trial of our experiment, the ResNet firstly scanned the input image to
188 identify and extract individual faces using a pre-trained face-detection model integrated
189 in the `face_recognition` library. The image of each face was then mapped to a
190 128-dimensional vector space where images of the same person are likely to be near to
191 each other (i.e., face encoding). Then, the ResNet computed the difference between the
192 encoding of each face and the encoding of the target face, normalised using the
193 Frobenius norm. If any extracted face had a difference smaller than a threshold, t , the
194 stimulus was labelled as “target”, otherwise it was labelled as “nontarget”. The
195 threshold was set to the value that maximised the accuracy on the training set using
196 8-fold cross-validation (average threshold across the eight folds, $t = 0.526 \pm 0.006$).
197

In each trial, the confidence weight of the ResNet, w_{ResNet} , was computed as follows:
198

$$w_{ResNet} = \begin{cases} \frac{2}{1+e^{-10 \cdot \frac{t-d}{t}}} - 1, & \text{if } d < t \text{ (“target”)} \\ \frac{2}{1+e^{-10 \cdot \frac{d-t}{1-t}}} - 1, & \text{if } d \geq t \text{ (“nontarget”)} \end{cases}, \quad (1)$$

where d is the difference between the encodings of the face extracted from the image
199 and the target face (normalised using the Frobenius norm). If more than one face was
200 extracted from an image, only the face with the minimum d was considered. Fig 2
201 shows a plot of the function in Equation (1) for the average t . This formula ensures that
202 the ResNet has substantially bigger confidence weights when the difference between the
203

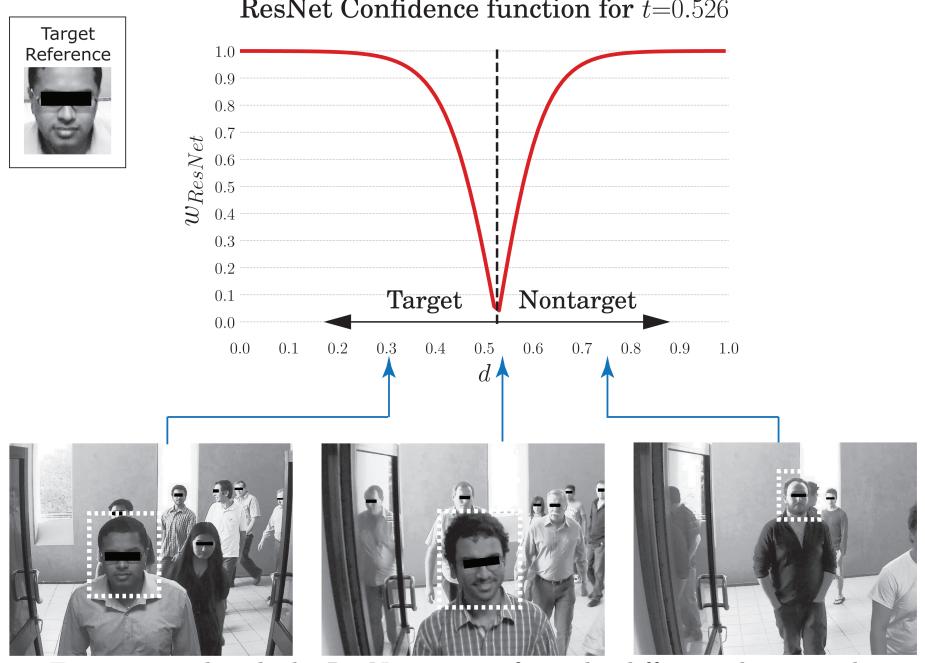


Fig 2. Function used with the ResNet to transform the difference between the encodings of the stimulus and target faces, d , into confidence weights for the average value of the threshold t over the eight folds. Examples of stimuli with $d=0.301$ (left), $d=0.535$ (middle), and $d=0.750$ (right) are also shown at the bottom. In each image, the bounding box of the face with the minimum d is highlighted. The reference image of the target person is shown on the top left. Images representing people have been blurred for publication requirements.

target face and the stimulus is very small or very big (confident decisions) than when it
is close to the threshold (not confident decisions).

Trials where the ResNet did not identify any face in the image were labelled as
“nontarget”, given the prevalence of nontarget stimuli in the training set (i.e., 75%). In
such trials, the corresponding weights w_{ResNet} were set to 0, so that, in cyborg groups,
the ResNet vote would be ignored and only the humans influenced the decision.

Making group decisions.

Groups of size $m = 2, \dots, 10$ were formed off-line by combining the $N=10$ participants
in all possible $(N)_m$ ways. This procedure led to have 45 groups of size 2, 120 groups of
size 3, etc. Group decisions were then obtained using different methods. These are
illustrated in Fig 3 and are described in detail below.

Traditional and BCI-assisted groups were composed only by human participants
(Fig 3(a) and (b), respectively). In this case, the decision of a group of size m in a trial

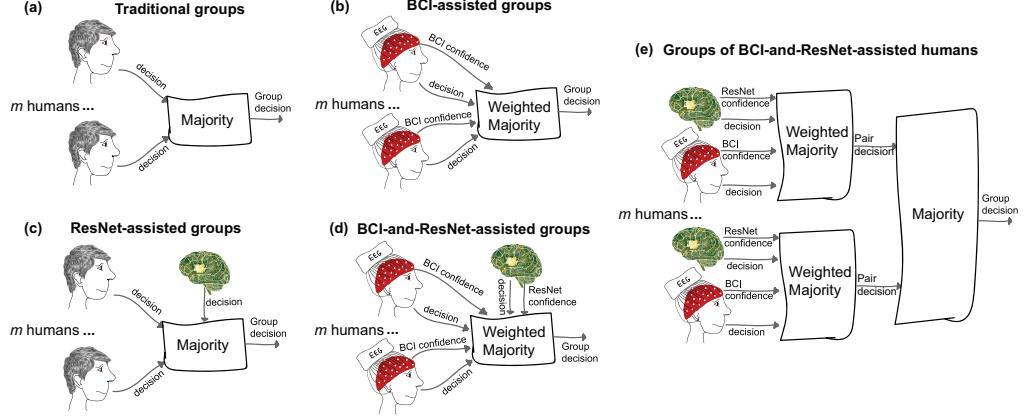


Fig 3. Summary of the different methods for obtaining decisions of groups of size m . (a) Traditional groups of human participants based on standard majority. (b) Groups of human participants weighing individual decisions with the confidence estimated by the BCI. (c) Traditional groups based on standard majority with an additional member being the ResNet. (d) BCI-assisted groups with the ResNet as an additional member, with its degree of confidence. (e) Groups where each participant is assisted by the BCI and paired with the ResNet, and the decisions of these pairs are then integrated using standard majority.

was given by:

$$d_{group} = \text{rsign} \left(\sum_{p=1}^m w_p \cdot d_p \right),$$

where $d_p = -1$ if user p decided the trial represented a “target” and $d_p = +1$ for a “nontarget”, w_p is the corresponding weight, and rsign a randomising sign operator which returns $+1$ (“nontarget”) if its argument is positive, -1 (“target”) if it is negative, and randomly chooses between $+1$ and -1 if its argument is 0. Traditional group decisions were obtained using standard majority, i.e., $w_p=1$ for all trials and users, while BCI-assisted group decisions used the decision confidence estimated by the BCI from the brain signals as weight w_p .

ResNet-assisted and BCI-and-ResNet-assisted groups (Fig 3(c) and (d), respectively) were composed by m human participants plus the ResNet. Their decisions in a trial were given by:

$$d_{ResNet_group} = \text{rsign} \left(\sum_{p=1}^m w_p \cdot d_p + w_{ResNet} \cdot d_{ResNet} \right),$$

where $d_{ResNet} \in \{-1, 1\}$ is the decision of the ResNet in the trial, and w_{ResNet} is the corresponding weight. ResNet-assisted group decisions were obtained by using standard majority, i.e., $w_p=w_{ResNet}=1$ for all trials and participants, while

BCI-and-ResNet-assisted group decisions were obtained by using the decision confidence
 estimated by the BCI from the brain signals as weight w_p , and the confidence of the
 ResNet w_{ResNet} as weight for its decision.

Finally, groups of m BCI-and-ResNet-assisted humans (Fig 3(e)) were obtained by
 firstly pairing each participant with the ResNet and computing their d_{ResNet_group}
 decisions. These m decisions were then integrated using standard majority to obtain the
 decision of a group of BCI-and-ResNet-assisted humans:

$$d_{ResNet_humans_group} = \text{rsign} \left(\sum_{p=1}^m \text{rsign}(w_p \cdot d_p + w_{ResNet} \cdot d_{ResNet}) \right).$$

This last approach was suggested by literature for human groups [47] showing that
 aggregating votes from smaller inner crowds outperforms pooling votes from larger
 groups of individual decision makers.

Results

The ResNet performs better than the average individual, but with lower sensitivity.

The accuracy of the participants in the experiment was quite varied, ranging from
 52.8% to 92.4%, with an average accuracy of 72.3% and a standard deviation (SD) of
 12.0%. We have also looked at the ability of each participant of identifying nontargets
 and targets by computing their specificity and sensitivity as:

$$\begin{aligned} \text{Specificity} &= \frac{\text{True Negatives}}{\text{True Negatives} + \text{False Positives}}, \\ \text{Sensitivity} &= \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}. \end{aligned}$$

The participants' average specificity was 77.4% (SD=15.8%) and their average
 sensitivity was 56.9% (SD=10.9%). Hence, they recognised nontarget stimuli much
 better than target ones.

The ResNet had an average accuracy of 84.4%, a specificity of 98.6% and a

sensitivity of 41.7%. Hence, it was substantially more accurate than the average participant. However, the ResNet was particularly good in assessing nontarget images (i.e., high specificity), but quite weak in detecting targets (i.e., low sensitivity) compared to the average participant. Since target stimuli appeared only in 25% of the trials, the ResNet could still achieve high accuracy, which was the metrics used to choose the threshold t (see *Methods*). In two nontarget trials and one target trial, the ResNet did not identify any face in the image and, therefore, classified the picture as “nontarget” (see *Methods*). Hence, only one error made by the ResNet was caused by its first step (i.e., extracting faces), which we could assume to be a trivial task for humans.

One may wonder whether a ResNet using a different metrics to optimise the choice of the threshold t could achieve better sensitivity. When choosing t in order to maximise the F1 score in cross-validation instead of the accuracy, the ResNet reduced its accuracy (82.6%) and specificity (94.0%), and increased its sensitivity (48.6%). The average value of the threshold t over the eight folds was 0.559, while it was 0.526 when maximising accuracy. Similar results were obtained when maximising Cohen’s kappa to optimise t (average $t=0.543$, accuracy=81.6%, specificity=94.0%, and sensitivity=44.4%). In all cases, the sensitivity of the ResNet was worse than that of the average participant.

The ResNet distinguishes itself from the crowd.

To exhibit error-correction capability, groups require their members to make errors in different trials [48]. In fact, if all group members make the same, incorrect decision in a trial, the group will also make an incorrect decision. To assess error-correction capability of pairs, we used the Hamming loss, i.e., the fraction of trials in which the decisions of the two users were different. We computed the Hamming loss of each pair of participants, as well as the Hamming loss of each user when paired with the ResNet. We also computed the Hamming loss between the decisions of each user or the ResNet and the ideal user (i.e., one that is always correct), which corresponds to their error rates.

Participants had performance better than random and, therefore, their decisions were the same in the majority of the trials. However, the Hamming loss between two participants ranged from 12.5% and 47.2% (mean=30.67%, SD=7.92%). This suggests that the pairs exhibited error-correction capability in a large percentage of trials.

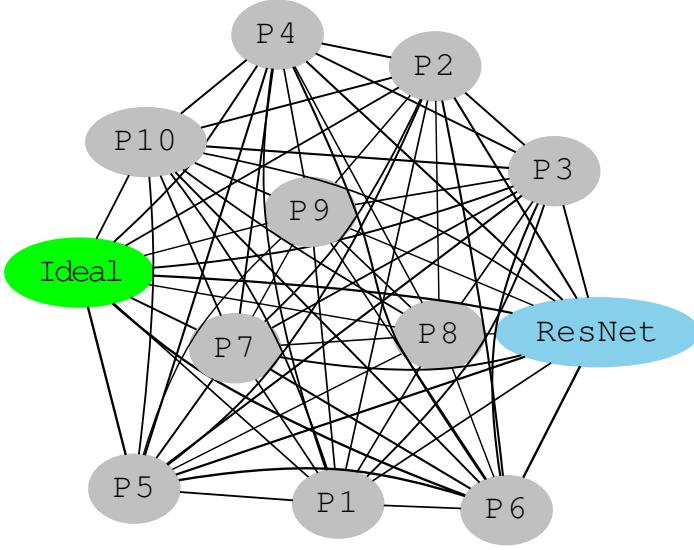


Fig 4. Graphical representation of the Hamming loss between the decisions of human participants (grey nodes), ResNet (light blue node) and ideal decision maker (green node). The weight of each edge is inversely proportional to the Hamming loss between the two nodes connected by that edge, hence representing how similar were the decisions made by the two nodes.

When looking at the Hamming loss between the ResNet and each participant, we
283 found similar results. In this case, the Hamming loss ranged between 12.2% and 48.3%
284 (mean=28.92%, SD=11.60%). A Wilcoxon signed-rank test showed no differences
285 between the Hamming loss of humans and ResNet ($p=0.646$). Given its high
286 performance, of course the ResNet had minimum Hamming loss with the participants
287 with the highest accuracy (i.e., P5 and P6). Yet, even when paired with them, the pair
288 still exhibited error-correction capability in more than 12% of the trials.
289

Different decision makers can be seen as nodes in a weighted graph, with the
290 Hamming loss measure between pairs of decision makers representing the weight
291 (dissimilarity) between corresponding nodes. Then, utilising a graph layout algorithm
292 one can flatten the graph to obtain an intuition of similarities and differences in
293 decision-making behaviour. Fig 4 shows the result of this process when using the
294 program `neato` from the **Graphviz** library.
295

This diagram as well as the previous analyses show how the ResNet behaves
296 differently from virtually all other participants.
297

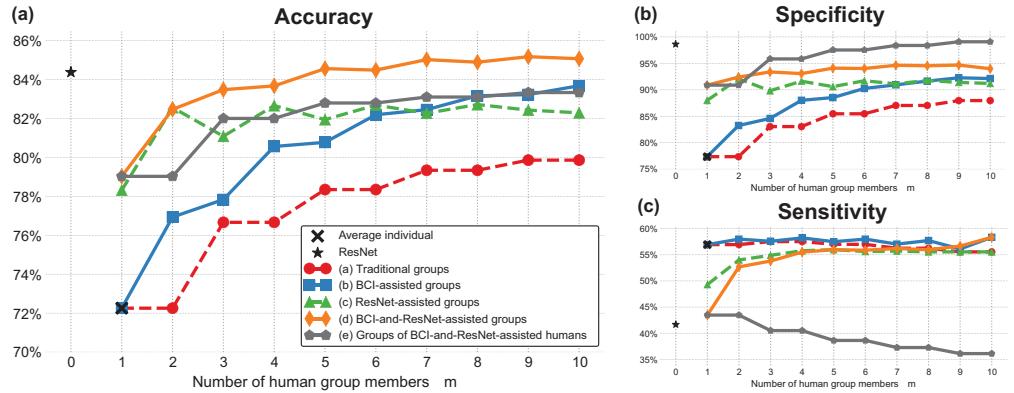


Fig 5. Average performance of individuals, ResNet, and groups of increasing size making decisions using the approaches described in Fig 3. **(a)** Accuracy. **(b)** Specificity (true negative rate). **(c)** Sensitivity (true positive rate).

Cyborgs have significantly higher performance than purely human groups.

Fig 5 reports the average accuracy, specificity and sensitivity of groups of increasing number of human members (m) using the five different methods for making group decisions analysed in this study (see Fig 3). The average performance of individuals and of the ResNet are also reported for reference.

Let us first focus on the accuracy (Fig 5(a)). As one would expect [16], the more opinions are integrated in the group decisions, the more accurate the group becomes. However, it is also apparent from the figure that the method chosen to integrate individual opinions and obtain group decisions is at least as important in determining group performance as the group size. In particular, it is clear that traditional groups of humans (red line in Fig 5(a)) using standard majority are outperformed by all other four methods, regardless the group size.

Of course, one may expect methods which use some form of confidence estimation on a decision-by-decision basis to break ties, such as the BCI-assisted groups (blue line in Fig 5(a)), to be superior to methods using simple majority and breaking ties randomly. However, the performance of ResNet-assisted groups (green line in Fig 5(a)), which do not make use of confidence estimates, *appears* to contradict this expectation. Such groups are not only more accurate than equally-sized traditional groups, but also more accurate than BCI-assisted groups of sizes 2–6. This happens because the ResNet acts as an additional, very accurate member within the group (the ResNet is not only

significantly better than the average participant, but only second to P5 and P6 in terms
319 of accuracy). Since ResNet-assisted groups are based on standard majority, the smaller
320 the group, the bigger the influence of the ResNet on the group decision. Indeed, the
321 biggest improvement brought by ResNet-assisted groups over traditional groups occurs
322 for $m=2$, where the addition of the ResNet allows us to turn most of the ties generated
323 by the two humans into correct decisions. This results in a boost in accuracy of more
324 than 10% over traditional pairs. For bigger group sizes, the difference in average
325 performance between ResNet-assisted and traditional groups gradually diminishes. We
326 should expect the performance of ResNet-assisted groups to eventually converge to that
327 of traditional groups. Yet, even for the larger group sizes considered in our study,
328 ResNet-assisted groups were still better than traditional groups by 2–3%.
329

Given the benefits provided by assisting groups with the BCI and the ResNet
330 individually, it is not surprising to see that assisting groups with both (orange line in
331 Fig 5(a)) yields even better performance — in fact, the best performance over all other
332 methods of obtaining group decisions.
333

Interestingly, when using groups of BCI-and-ResNet-assisted humans (grey line in
334 Fig 5(a)), groups were generally more accurate than ResNet-assisted groups (the only
335 exceptions being groups of size $m=2$ and $m=4$), but less accurate than
336 BCI-and-ResNet-assisted groups. This indicates that, contrary to what is suggested in
337 the literature for human groups [47], aggregating votes from smaller inner crowds does
338 not outperform pooling votes from larger groups of individual decision makers, when
339 crowds are composed by humans and artificial agents.
340

To statistically validate the aforementioned observations, we used the Wilcoxon
341 signed-rank test to perform pairwise comparisons of the accuracy of groups of increasing
342 sizes using the five methods analysed in the study. Table 1 reports the two-sided
343 p -values returned by the test. Statistical analyses for groups of size 10 was not possible
344 as we only had one such group.
345

As expected, BCI-assisted groups are significantly more accurate than equally-sized
346 traditional groups based on standard majority for all group sizes – see “(a) vs (b)”
347 column in Table 1. This confirms previous results obtained with this collaborative BCI
348 both with the same set up (but using a combination of different neural features and
349 reaction times) [41], and in other visual-search experiments [40]. Moreover,
350

Table 1. Two-sided p -values of the Wilcoxon signed-rank test comparing the performance of: (a) traditional groups, (b) BCI-assisted groups, (c) ResNet-assisted groups, (d) BCI-and-ResNet-assisted groups, and (e) groups of BCI-and-ResNet-assisted humans. Values in bold represent $p < 0.05$ confidence level.

Group size (# groups)	(a) vs (b)	(a) vs (c)	(b) vs (c)	(c) vs (d)	(b) vs (d)	(d) vs (e)
1 ($n = 10$)	—	1.7×10^{-2}	1.7×10^{-2}	8.8×10^{-1}	2.8×10^{-2}	—
2 ($n = 45$)	1.1×10^{-7}	7.6×10^{-9}	3.7×10^{-8}	3.2×10^{-1}	2.5×10^{-8}	8.5×10^{-6}
3 ($n = 120$)	6.0×10^{-15}	2.5×10^{-20}	1.2×10^{-19}	5.1×10^{-18}	1.2×10^{-20}	2.1×10^{-7}
4 ($n = 210$)	1.5×10^{-35}	3.3×10^{-36}	2.8×10^{-28}	6.7×10^{-21}	7.9×10^{-33}	2.5×10^{-13}
5 ($n = 252$)	2.4×10^{-39}	9.2×10^{-43}	7.7×10^{-27}	2.4×10^{-42}	4.7×10^{-43}	4.0×10^{-17}
6 ($n = 210$)	3.3×10^{-36}	3.2×10^{-36}	4.2×10^{-8}	2.2×10^{-32}	3.2×10^{-34}	5.0×10^{-19}
7 ($n = 120$)	2.8×10^{-21}	1.9×10^{-21}	5.3×10^{-2}	2.8×10^{-21}	2.6×10^{-21}	8.2×10^{-14}
8 ($n = 45$)	5.1×10^{-9}	5.1×10^{-9}	8.4×10^{-4}	5.1×10^{-9}	7.4×10^{-9}	2.7×10^{-7}
9 ($n = 10$)	5.0×10^{-3}	5.1×10^{-3}	9.3×10^{-3}	5.0×10^{-3}	4.9×10^{-3}	9.3×10^{-3}

ResNet-assisted groups were also significantly more accurate than traditional groups of 351 only humans for all group sizes, as per “(a) vs (c)” column in Table 1. When comparing 352 ResNet-assisted and BCI-assisted groups (“(b) vs (c)” column in Table 1), the test 353 confirms that the former are significantly more accurate than the latter in small groups 354 (i.e., $m=2\text{--}6$). Conversely, BCI-assisted groups are significantly better than 355 ResNet-assisted groups for large group sizes ($m=8\text{--}9$). The two methods were on par for 356 groups of size 7. For all group sizes m , BCI-and-ResNet-assisted groups (orange line in 357 Fig 5) are significantly more accurate than groups assisted by only the ResNet or the 358 BCI (“(c) vs (d)” and “(b) vs (d)” columns of Table 1, respectively) for all group sizes 359 except for $m=2$, where BCI-and-ResNet-assisted groups and ResNet-assisted groups are 360 on par. BCI-and-ResNet-assisted groups are also significantly more accurate than 361 groups of BCI-and-ResNet-assisted humans (“(d) vs (e)” column in Table 1). 362

Let us now analyse the ability of groups of correctly detecting nontargets and targets 363 by looking at their specificity (Fig 5(b)) and sensitivity (Fig 5(c)). Here we should 364 remember that, as seen in the previous section, the ResNet achieved its high accuracy 365

thanks to the high specificity, while it suffered from low sensitivity. On the one hand, 366 similarly to what happened with accuracy, the specificity also increased with the group 367 size. Interestingly, in this case groups of BCI-and-ResNet-assisted humans (grey line) 368 have better specificity than equally-sized BCI-and-ResNet-assisted groups (orange line) 369 for $m = 3 - 10$, while the opposite happened for accuracy. On the other hand, the 370 sensitivity of groups of different sizes (Fig 5(c)) seemed to saturate around the average 371 human sensitivity (black cross). This suggests that participants found the target 372 difficult to spot in the same trials. To validate this hypothesis, we computed the 373 percentage of target trials in which more than half of the participants did not spot the 374 target. The majority of participants did not spot the target in 39% of the trials. Yet, 375 the average participants had a sensitivity better than random, while the ResNet did not. 376 As a result, when pairing the ResNet with a human participant to make groups of 377 BCI-and-ResNet-assisted humans (grey line), the sensitivity degraded with the increase 378 of the group size. Interestingly, when the ResNet was included in groups as an 379 additional, equal team member (orange and green lines in Fig 5(c)), groups were able to 380 improve their sensitivity as the number of human team members increased. These 381 results confirm that BCI-and-ResNet-assisted groups were our best method for making 382 group decisions. 383

BCI and ResNet confidence correlates with correctness. 384

Fig 6 shows the distribution of confidence values estimated by the BCI for the human 385 participants (left) and the distribution of confidence values of the ResNet (right), for 386 incorrect and correct decisions. The diagram on the left has been drawn with data from 387 799 incorrect trials and 2081 correct ones, while the diagram on the right has been 388 obtained from 45 incorrect trials and 243 correct ones. 389

Fig 6(left) indicates that, as we found in other experiments [40], the BCI is able to 390 assign significantly higher values of confidence to trials in which participants are correct 391 than to trials in which they are incorrect (Kruskal-Wallis test $p=1.6 \times 10^{-21}$). The 392 median BCI confidence values for incorrect and correct trials are 0.747 and 0.850, 393 respectively. The ResNet is also able to assess its degree of confidence via Equation 1 394 (see *Methods*) and its confidence values are lower in incorrect trials (median=0.793) 395

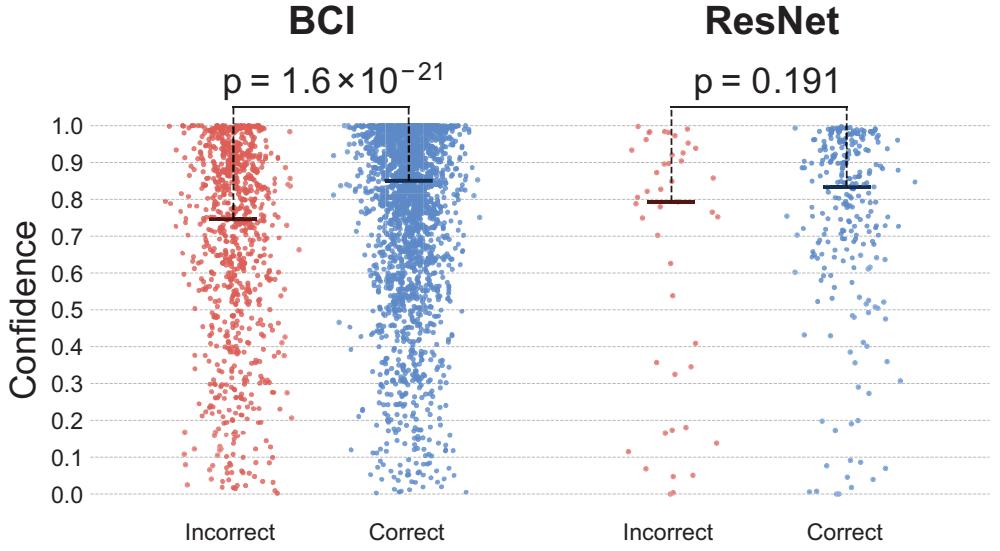


Fig 6. Distribution of confidence values for incorrect and correct decisions for the BCI (left) and the ResNet (right). Black lines indicate median values of each set. The p values of a Kruskal-Wallis test comparing the correct and incorrect distributions are also reported.

than in correct ones (median=0.834), although not significantly due to the smaller sample size (Kruskal-Wallis $p=0.191$). 396
397

From Fig 6 it is clear that, for both the BCI and the ResNet, incorrect trials have a much sparser distribution than correct ones. Indeed, the standard deviations of the BCI confidence for incorrect and correct trials are 0.268 and 0.239, respectively, while for the ResNet they are 0.342 and 0.238, respectively, and these differences are significant (Levene test $p=2.9 \times 10^{-7}$ and $p=2.0 \times 10^{-3}$, respectively). Once again, the difference in standard deviations between correct and incorrect confidence values is bigger for the ResNet. 398
399
400
401
402
403
404

Taken together, these statistical differences suggest that both the BCI and the ResNet are able to assess the decision confidence, albeit not perfectly. Of course, this is not entirely unexpected. The BCI bases its confidence estimates on the brain signals of the users, hence facing two major challenges that affect its certainty. Firstly, EEG signals are generally affected by a high level of noise caused by various sources (e.g., muscular activity, electrical interference, etc.). Secondly, humans participants may be overconfident or underconfident, two behaviours which might in turn also impact their brain patterns, thereby making it potentially difficult for our algorithms to reliably infer 405
406
407
408
409
410
411
412

the probability of a decision being correct. The ResNet decisions and confidence are 413
also affected by noise from different sources (e.g., partial occlusion of the target face). 414

Irrespective of the reasons for the differences in confidence distributions between 415
ResNet and BCI-assisted humans reported in Fig 6, due to such differences, in a 416
scenario where a BCI-assisted human works in pair with the ResNet and there is a tie, 417
the human is likely to break the tie in his/her favour. While this may or may not be 418
optimal from a group decision perspective, it may make working like a cyborg much 419
more acceptable to the human in the team. 420

Discussion 421

A cyborg is typically defined as the combination of a human and a machine, which is 422
either invasively embedded inside the human or replaces/extends a body part, with the 423
aim of restoring some function or augmenting human capabilities. These types of 424
cyborgs were once only the domain of science fiction, but are nowadays becoming a 425
reality thanks to the advancements in neural engineering and prosthetics. However, we 426
believe that the terms cyborg and cyborg group are also appropriate to describe 427
non-invasive combinations of one or more humans and one or more machines, such as 428
the systems described in this paper to augment performance in face recognition. 429

Of course, such systems can also be considered as forms of Human-Machine 430
Interaction (HMI), consisting of a machine analysing information, and humans using 431
this analysis to make informed decisions. In this area fall hybrid BCIs [49] that allow 432
users to control certain parts of an external device with their mind, while other parts 433
are controlled by another system. In this scenario, we generally refer to *shared-control* 434
systems, where the artificial component of the HMI generally takes care of all low-level 435
decisions [50]. For example, when driving a car [51] or controlling a wheelchair [52] with 436
a hybrid BCI, the autonomous control system usually performs the obstacle avoidance, 437
while the BCI, relying on noisy brain measures, only issues high-level commands [53]. 438
However, in those scenarios humans and machines do not fully interact: they simply 439
control different aspects of the device. 440

A more recent and promising avenue of research has been looking at Human-Machine 441
Cooperation (HMC), which further extend HMI to allow the human and the machine to 442

cooperate on a more equal footing and at a higher (cognitive) level [54]. This novel
443 approach could allow us to create a team of human and artificial operators with no a
444 priori dichotomy between high-level (human) and low-level (machine) elements. Instead,
445 the contribution of each operator to the control strategy or the decision-making process
446 depends on a time-dependent and situation-dependent assessment of each actor's
447 strengths and weaknesses. It is clear that in tasks which are challenging for both
448 humans and machines, joining forces may allow us to significantly improve performance.
449 However, it is not clear how the cooperation can be made to work practically.
450

The cyborgs and cyborg groups proposed in this study represent operational
451 solutions to the problem of how HMC could be implemented in an uncentralised and
452 self-organising manner in the domain of decision making. The simplest solution would
453 be creating groups where humans and machines have an equal weight in the decision
454 making (majority-based groups). In our face-recognition task, the ResNet achieved
455 significantly better performance than the average participant, hence allowing
456 ResNet-assisted groups to significantly improve performance over traditional groups of
457 humans. However, this approach is situation-independent, i.e., it is not able to
458 discriminate in which cases the machine or the human should be trusted more.
459

Cyborg group performance could be significantly boosted by assessing the decision
460 confidence of team members (computers or humans) in each decision and weighing their
461 decisions accordingly. We have shown that machine-vision algorithms could be designed
462 that also estimate their degree of confidence, while BCIs could be used to decode the
463 decision confidence of humans from their brain activity. These BCI-and-ResNet-assisted
464 groups fully implement the HMC strategy, allowing to decide the contribution to the
465 group decision of each team member (human or machine) on a decision-by-decision
466 basis, depending on the degree of confidence of each actor. Moreover, contrary to what
467 has been suggested in the literature for human groups [47], cyborg groups function
468 better when human and machine decisions are integrated at the same level, instead of
469 integrating decisions of human-machine pairs.
470

The cyborgs proposed in this study also take into account ethical and liability issues
471 concerning automated decision making. We have shown that the confidence values
472 provided by the BCI for humans are higher than the confidence of the ResNet, leading
473 to humans counting often more than the ResNet in uncertain decisions. When cyborg
474

groups are used at an operational level, it is ethically simpler (albeit not necessarily 475
optimal) to let humans decide in cases where there is no clear majority of opinions 476
between humans and ResNet. However, when humans in a group cannot agree on a 477
decision, it is more reasonable to trust fully-automated systems than flipping a coin to 478
break ties. 479

Overall, all strategies explored in this study on how to integrate humans and 480
machines have shown significant advantages in face-recognition performance over 481
traditional groups of only humans. When fully-automated machine-vision algorithms 482
are not reliable enough, BCIs could be used to significantly boost the performance of 483
groups of human operators. In other situations, pairing machines and BCI-assisted 484
humans may provide the best performance. Future work will verify these findings with 485
other tasks where automated algorithms could be developed. Furthermore, ensemble of 486
ResNets or similar artificial agents could be created and their performance compared 487
with that of groups of humans or humans and machines, to further verify the need of 488
humans in the loop to achieve maximum accuracy. 489

References 490

1. Eckstein MP. Visual search: A retrospective. *Journal of Vision*. 2011;11(5):1–36. 491
doi:10.1167/11.5.14. 492
2. Thorpe S, Fize D, Marlot C. Speed of processing in the human visual system. 493
Nature. 1996;381(6582):520–522. doi:10.1038/381520a0. 494
3. Nowakowska A, Clarke ADF, Hunt AR. Human visual search behaviour is far 495
from ideal. *Proceedings of the Royal Society of London B: Biological Sciences*. 496
2017;284(1849):1–6. doi:10.1098/rspb.2016.2767. 497
4. Eckstein MP, Koehler K, Welbourne LE, Akbas E. Humans, but Not Deep 498
Neural Networks, Often Miss Giant Targets in Scenes. *Current Biology*. 499
2017;27(18):2827–2832.e3. doi:10.1016/j.cub.2017.07.068. 500
5. Mack A, Rock I. Inattentional Blindness. Cambridge, MA: MIT Press; 1998. 501

6. Cinel C, Humphreys GW, Poli R. Cross-Modal Illusory Conjunctions Between Vision and Touch. *Journal of Experimental Psychology: Human Perception and Performance*. 2002;28(5):1243–1266. doi:10.1037/0096-1523.28.5.1243. 502
503
7. Robertson DJ, Burton M. Unfamiliar Face Recognition: Security, surveillance and smartphones. *The Journal of the Homeland Defense and Security Information Analysis Center*. 2016;3(1):14–21. 505
506
507
8. Robertson DJ, Jenkins R, Burton AM. Face detection dissociates from face identification. *Visual Cognition*. 2017;6285:1–9. 508
509
510
doi:10.1080/13506285.2017.1327465.
9. Ghuman AS, Brunet NM, Li Y, Konecky RO, Pyles JA, Walls SA, et al. 511
Dynamic encoding of face information in the human fusiform gyrus. *Nature Communications*. 2014;5(5672):5672. doi:10.1038/ncomms6672. 512
513
10. Robertson DJ, Noyes E, Dowsett AJ, Jenkins R, Burton AM. Face recognition by 514
metropolitan police super-recognisers. *PLoS ONE*. 2016;11(2):1–8. 515
516
doi:10.1371/journal.pone.0150036.
11. Zhao W, Chellappa R, Phillips PJ, Rosenfeld A. Face Recognition: A Literature 517
Survey. *ACM Computing Survey*. 2003;35(4):399–458. 518
12. Yang MH, Kriegman DJ, Ahuja N. Detecting faces in images: A survey. *IEEE Transactions on Pattern Analysis and Machine Intelligence*. 2002;24(1):34–58. 519
520
521
doi:10.1109/34.982883.
13. Sinha P, Balas B, Ostrovsky Y, Russell R. Face recognition by humans: Nineteen 522
results all computer vision researchers should know about. *Proceedings of the IEEE*. 2006;94(11):1948–1961. doi:10.1109/JPROC.2006.884093. 523
524
14. Tan X, Chen S, Zhou ZH, Zhang F. Face recognition from a single image per 525
person: A survey. *Pattern Recognition*. 2006;39(9):1725–1745. 526
527
doi:10.1016/j.patcog.2006.03.013.
15. Burton AM, Wilson S, Cowan M, Bruce V. Face Recognition in Poor-Quality 528
Video: Evidence From Security Surveillance. *Psychological Science*. 1999;10(3):243–248. doi:10.1111/1467-9280.00144. 529
530

16. Bang D, Frith CD. Making better decisions in groups. Royal Society Open Science. 2017;4(8):170193. doi:10.1098/rsos.170193. 531
532
17. Surowiecki J. The Wisdom of Crowds. New York: Anchor; 2005. 533
18. Gürçay B, Mellers BA, Baron J. The Power of Social Influence on Estimation Accuracy. Journal of Behavioral Decision Making. 2014;28:250–261. 534
535
doi:10.1002/bdm.1843. 536
19. Kurvers RHJM, Herzog SM, Hertwig R, Krause J, Carney PA, Bogart A, et al. Boosting medical diagnostics by pooling independent judgments. Proceedings of the National Academy of Sciences. 2016;113(31):8777–8782. 537
538
doi:10.1073/pnas.1601827113. 539
540
20. White D, Burton AM, Kemp RI, Jenkins R. Crowd effects in unfamiliar face matching. Applied Cognitive Psychology. 2013;27(6):769–777. 541
542
doi:10.1002/acp.2971. 543
21. Kerr NL, Tindale RS. Group Performance and Decision Making. Annual Review of Psychology. 2004;55(1):623–655. doi:10.1146/annurev.psych.55.090902.142009. 544
545
22. Branson L, Steele NL, Sung CH. When two heads are worse than one: Impact of group style and information type on performance evaluation. Journal of Business and Behavioral Sciences. 2010;22(1):75–84. 546
547
548
23. Lorenz J, Rauhut H, Schweitzer F, Helbing D. How social influence can undermine the wisdom of crowd effect. Proceedings of the National Academy of Sciences. 2011;108(22):9020–9025. doi:10.1073/pnas.1008636108. 549
550
551
24. Baharad E, Goldberger J, Koppell M, Nitzan S. Distilling the wisdom of crowds: Weighted aggregation of decisions on multiple issues. Autonomous Agents and Multi-Agent Systems. 2011;22(1):31–42. doi:10.1007/s10458-009-9120-y. 552
553
554
25. Laan A, Madirolas G, de Polavieja GG. Rescuing Collective Wisdom when the Average Group Opinion Is Wrong. Frontiers in Robotics and AI. 2017;4(November). doi:10.3389/frobt.2017.00056. 555
556
557

26. Marshall JAR, Brown G, Radford AN. Individual Confidence-Weighting and Group Decision-Making. *Trends in Ecology and Evolution*. 2017;32(9):636–645. doi:10.1016/j.tree.2017.06.004. 558
559
560
27. Pouget A, Drugowitsch J, Kepcs A. Confidence and certainty: distinct probabilistic quantities for different goals. *Nature Neuroscience*. 2016;19(3):366–374. doi:10.1038/nn.4240. 561
562
563
28. Navajas J, Bahrami B, Latham PE. Post-decisional accounts of biases in confidence. *Current Opinion in Behavioral Sciences*. 2016;11:55–60. 564
565
566
29. Sanchez C, Dunning D. Overconfidence among beginners: Is a little learning a dangerous thing? *Journal of Personality and Social Psychology*. 2018;114(1):10–28. doi:10.1037/pspa0000102. 567
568
569
30. Bang D, Aitchison L, Moran R, Herce Castanon S, Rafiee B, Mahmoodi A, et al. Confidence matching in group decision-making. *Nature Human Behaviour*. 2017;1(May):1–7. doi:10.1038/s41562-017-0117. 570
571
572
31. Boldt A, Yeung N. Shared Neural Markers of Decision Confidence and Error Detection. *Journal of Neuroscience*. 2015;35(8):3478–3484. 573
574
575
32. Selimbeyoglu A, Keskin-ergen Y, Demiralp T. Clinical Neurophysiology What if you are not sure? Electroencephalographic correlates of subjective confidence level about a decision. *Clinical Neurophysiology*. 2012;123(6):1158–1167. 576
577
578
579
33. Luck SJ, Hillyard SA. Electrophysiological correlates of feature analysis during visual search. *Psychophysiology*. 1994;31(3):291–308. 580
581
582
34. van den Berg B, Appelbaum LG, Clark K, Lorist MM, Woldorff MG. Visual search performance is predicted by both prestimulus and poststimulus electrical brain activity. *Scientific Reports*. 2016;6:37718. doi:10.1038/srep37718. 583
584
585

35. Luce RD. Response Times: Their Role in Inferring Elementary Mental Organization. vol. 8. Oxford University Press; 1986. 586
36. Kiani R, Corthell L, Shadlen MN. Choice Certainty Is Informed by Both Evidence and Decision Time. *Neuron*. 2014;84(6):1329–1342. 588
doi:10.1016/j.neuron.2014.12.015. 589
37. Ratcliff R, Sterns JJ. Modeling confidence judgments, response times, and multiple choices in decision making: Recognition memory and motion discrimination. *Psychological Review*. 2013;120(3):697–719. 591
doi:10.1037/a0033152. 592
38. Poli R, Valeriani D, Cinel C. Collaborative Brain-Computer Interface for Aiding Decision-Making. *PLoS ONE*. 2014;9(7):e102693. 595
doi:10.1371/journal.pone.0102693. 596
39. Valeriani D, Poli R, Cinel C. Enhancement of Group Perception via a Collaborative Brain-Computer Interface. *IEEE Transactions on Biomedical Engineering*. 2017;64(6):1238–1248. doi:10.1109/TBME.2016.2598875. 598
600
40. Valeriani D, Cinel C, Poli R. Group Augmentation in Realistic Visual-Search Decisions via a Hybrid Brain-Computer Interface. *Scientific Reports*. 2017;7(1):7772. doi:10.1038/s41598-017-08265-7. 601
603
41. Valeriani D, Cinel C, Poli R. Augmenting group performance in target-face recognition via collaborative brain-computer interfaces for surveillance applications. In: 2017 8th International IEEE/EMBS Conference on Neural Engineering (NER). Shanghai, China: IEEE; 2017. p. 415–418. Available from: <http://ieeexplore.ieee.org/document/8008378/>. 604
605
606
607
608
42. Phillips PJ, Yates AN, Hu Y, Hahn CA, Noyes E, Jackson K, et al. Face recognition accuracy of forensic examiners, superrecognizers, and face recognition algorithms. *Proceedings of the National Academy of Sciences*. 2018;115(24):6171–6176. doi:10.1073/pnas.1721355115. 609
610
611
612
43. Wong Y, Chen S, Mau S, Sanderson C, Lovell BC. Patch-based probabilistic image quality assessment for face selection and improved video-based face 613
- 614

- recognition. In: CVPR 2011 WORKSHOPS. IEEE; 2011. p. 74–81. Available from: <http://ieeexplore.ieee.org/document/5981881/>. 615
44. Ramoser H, Muller-Gerking J, Pfurtscheller G. Optimal Spatial Filtering of Single Trial EEG During Imagined Hand Movement. IEEE Transactions on Rehabilitation Engineering. 2000;8(4):441–6. 617
618
619
45. He K, Zhang X, Ren S, Sun J. Deep Residual Learning for Image Recognition. In: 2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR). IEEE; 2016. p. 770–778. Available from: 620
621
622
623
<http://ieeexplore.ieee.org/document/7780459/>.
46. Huang GB, Ramesh M, Berg T, Learned-Miller E. Labeled faces in the wild: A database for studying face recognition in unconstrained environments; 2007. 624
625
47. Navajas J, Niella T, Garbulsky G, Bahrami B, Sigman M. Aggregated knowledge from a small number of debates outperforms the wisdom of large crowds. Nature Human Behaviour. 2018; p. 1–27. doi:10.1038/s41562-017-0273-4. 626
627
628
48. Eckstein MP, Das K, Pham BT, Peterson MF, Abbey CK. Neural decoding of collective wisdom with multi-brain computing. NeuroImage. 2012;59(1):94–108. doi:10.1016/j.neuroimage.2011.07.009. 629
630
631
49. Pfurtscheller G, Allison BZ, Brunner C, Bauernfeind G, Solis-Escalante T, Scherer R, et al. The hybrid BCI. Frontiers in Neuroscience. 2010;doi:10.3389/fnpro.2010.00003. 632
633
634
50. Müller-Putz GR, Breitwieser C, Cincotti F, Leeb R, Schreuder M, Leotta F, et al. Tools for brain-computer interaction: a general concept for a hybrid BCI. Frontiers in Neuroinformatics. 2011;5. doi:10.3389/fninf.2011.00030. 635
636
637
51. Göhring D, Latotzky D, Wang M, Rojas R. Semi-autonomous Car Control Using Brain Computer Interfaces. In: Proceedings of the 12th International Conference on Intelligent Autonomous Systems (IAS). Springer Berlin Heidelberg; 2013. p. 393–408. 638
639
640
641

52. Philips J, Millan JDR, Vanacker G, Lew E, Galan F, Ferrez PW, et al. Adaptive shared control of brain actuated simulated wheelchair,. In: IEEE 10th International Conference on Rehabilitation Robotics; 2007. p. 408–414.
53. Tonin L, Leeb R, Tavella M. The role of shared-control in BCI-based telepresence. In: IEEE International Conference on Systems, Man, and Cybernetics. Istanbul; 2010. p. 1462–1466. Available from: <http://ieeexplore.ieee.org/document/5642338/>.
54. Hoc JM. From human – machine interaction to human – machine cooperation. Ergonomics. 2000;43(7):833–843. doi:10.1080/001401300409044.

Acknowledgements

The Authors acknowledge support of the UK Defence Science and Technology Laboratory (Dstl) and Engineering and Physical Research Council (EPSRC) under grant EP/P009204/1. This is part of the collaboration between US DOD, UK MOD and UK EPSRC under the Multidisciplinary University Research Initiative. The Authors also acknowledge support by the Defence and Security PhD programme through Dstl. Finally, the authors acknowledge support of the UK Dstl (DSTLX 1000128890) and US DOD (W911NF1810434) under the Bilateral Academic Research Initiative programme.

Author contributions statement

D.V. and R.P. conceived and designed the study. D.V. collected and analysed the data. D.V. and R.P. interpreted the data, wrote and revised the manuscript.